Lecture 8: EDA

CS109A Introduction to Data Science

Pavlos Protopapas, Kevin Rader and Chris Tanner

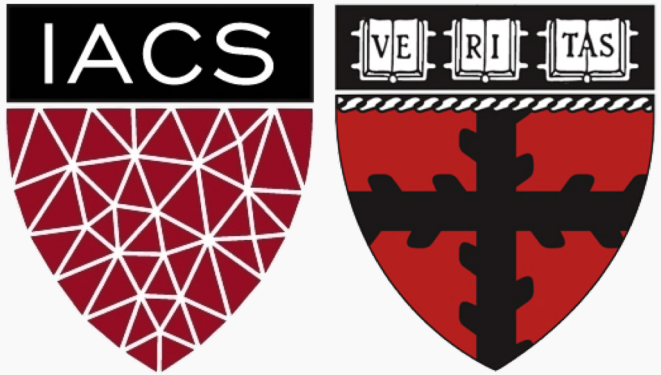

Lecture Outline

Data Science Process Example

- Dataset considerations
  - Comprehensive vs Sampled
  - Biases

Visualization

- Exploration (EDA)
- Communication

Lecture Outline

Data Science Process Example

- Dataset considerations
  - Comprehensive vs Sampled
  - Biases

Visualization

- Exploration (EDA)
- Communication

Example

Let's say that we are interested in the English Premier League (football/soccer) and want to build a model to predict a player's market value.

Question Does age affect one's market value?

Example

What do we do?

Example

What do we do?

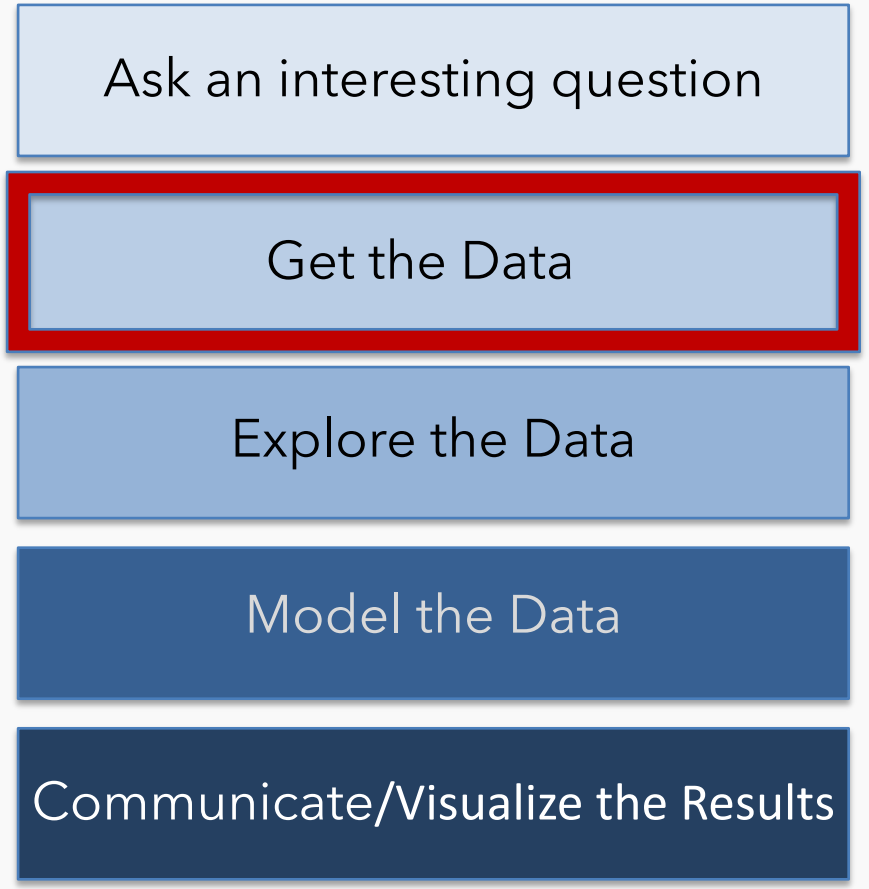

Dataset Considerations

- What data is necessary to answer our question?
- Is the source credible/authoritative? (.com, .net, .org, .gov, .name)
  company
  organization
- How difficult is it to analyze the dataset? (photos, videos, text?)
- What is the allowed usage of data under its license?
- Who collected the data?
- When was the data collected?

Dataset Considerations (continued)

- How was the data collected?
- How is the data formatted?
- Confidentiality concerns
- Does your data collection procedures need to be approved by an IRB?
- Comprehensive data vs sampled data?
- Biases

Dataset Considerations (continued)

- How was the data collected?
- How is the data formatted?
- Confidentiality concerns
- Does your data collection procedures need to be approved by an IRB?
- Comprehensive data vs sampled data?
- Biases

Lecture Outline

Data Science Process Example

- Dataset considerations
  - Comprehensive vs Sampled
  - Biases

Visualization

- Exploration (EDA)
- Communication

Dataset Considerations: Comprehensive Data

- We have access to all the data points that exist, which is usually a lot
- Collected and digitized as part of generalized procedures of an institution

The New York Times

13 million articles

~500 million tweets per day

CONGRESS.GOV

100,000s votes per year

Dataset Considerations: Sampled Data

- When collecting individual data is relatively expensive
- Only a portion of the population is sampled
- Not just restricted to polling or surveys

GALLUP

IMDb

1. Clover Food Lab

821 reviews

$$ · American (New), Sandwiches, Cafes

THE Q EVALUATIONS

nielsen

Lecture Outline

Data Science Process Example

- Dataset considerations
  - Comprehensive vs Sampled
  - Biases

Visualization

- Exploration (EDA)
- Communication

Dataset Considerations: Biases

- **A bias in sampled data occurs when a procedure causes the sample to overrepresent a subpopulation**
- Biases may not necessarily be intentional
- Even if you don't think over-representation of a subpopulation will bias the dataset with regard to your question, it's still a bias
- Always strive to minimize any biases in your data collection procedures

Dataset Considerations: Biases

Gallup Polls

- Randomly calls two groups of ~500 people a day by sampling among all possible phone numbers
- For landlines, asks for household member who has the next birthday
- Calls people living in all 50 states
- Tries to assure 70% cellphone, 30% landlines
- Weights data to reflect the demographics of the general population

Dataset Considerations: Biases

IMDb Movie Ratings

- Registered users rate films 1-10 stars; they are an overrepresented subpopulation relative to the general population
- Registered users who rate movies in their free time further over represents a specific segment of the general population
- *"Men Are Sabotaging The Online Reviews Of TV Shows Aimed At Women*[1]*"*
    - 60% who rated Sex in the City were women. Women gave it a 8.1, men gave it 5.8.

[1] **fivethirtyeight.com**

Dataset Considerations: Biases

IMDb Movie Ratings

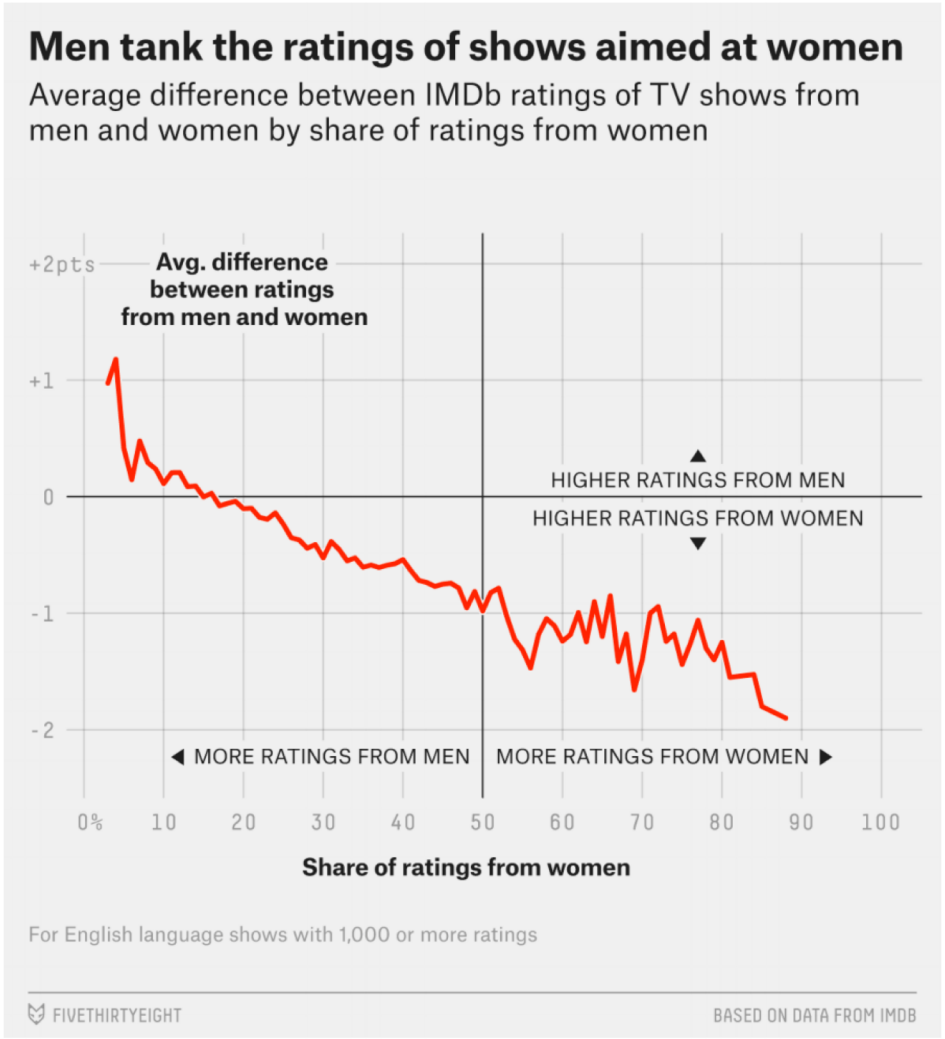

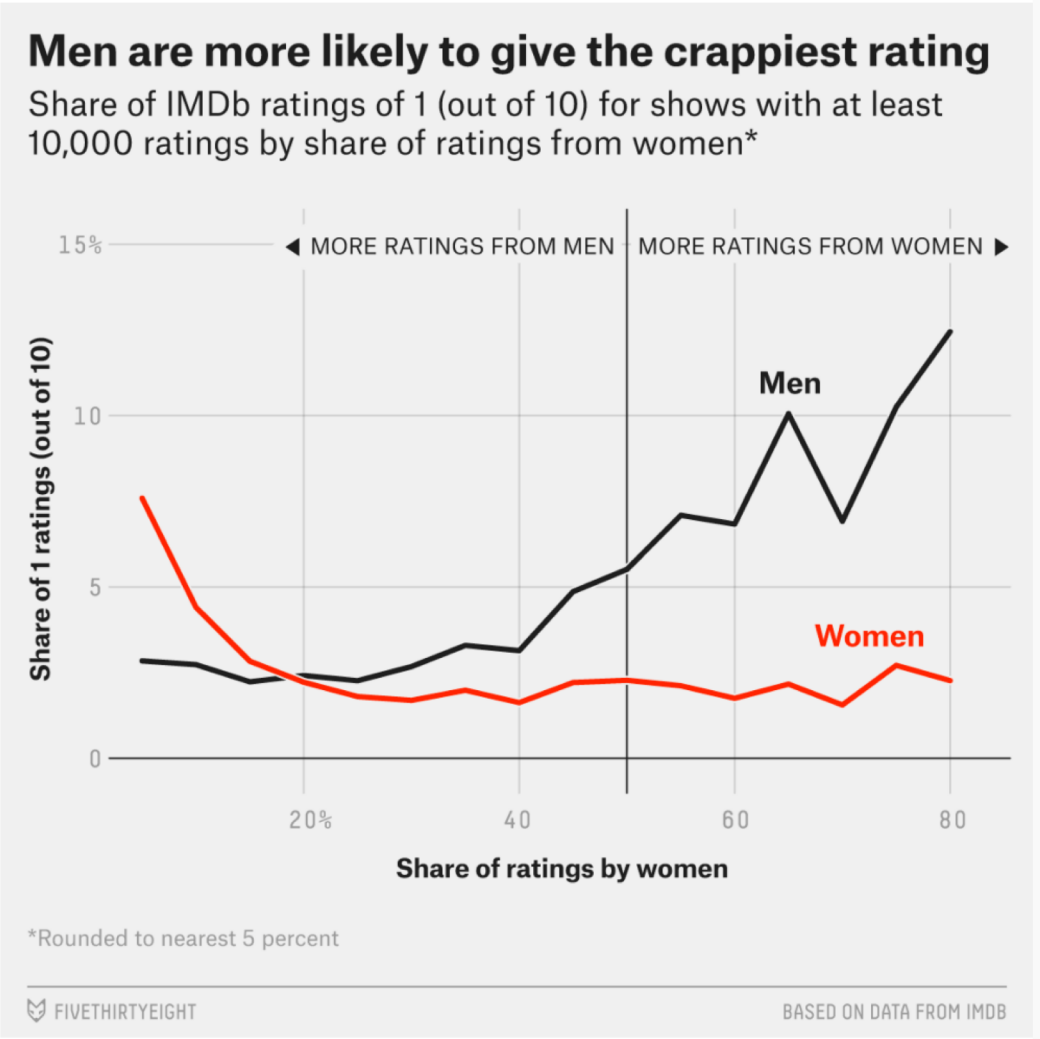

Dataset Considerations: Biases

Yelp Reviews

- Registered users rate businesses on a 1-5 star scale
- Registered users tend to represent a certain subset of the population (those who are more social media inclined and opinionated)
- Customers with extreme experiences are more likely to voice their opinions

Dataset Considerations: Biases

Yelp Reviews

6. Clover Food Lab

104 reviews

$ · Sandwiches, Cafes, American (New)

1. Clover Food Lab

821 reviews

$$ · American (New), Sandwiches, Cafes

Dataset Considerations: Biases

Yelp Reviews

6. Clover Food Lab

104 reviews

$ · Sandwiches, Cafes, American (New)

1. Clover Food Lab

821 reviews

$$ · American (New), Sandwiches, Cafes

Longwood Medical

Harvard Square

The actual price is the same but the rating shows difference. The population making the rating is different. Longwood - doctors, HS - tourists.

Back to our example…

Let's say that we are interested in the English Premier League (football/soccer) and want to build a model to predict a player's market value.

Question Does age affect one's market value?

Example: Get the data

name	club	age	position	market value
Alexis Sanchez	Arsenal	28	LW	65
Mesut Ozil	Arsenal	28	AM	50
Petr Cech	Arsenal	35	GK	7
Theo Walcott	Arsenal	28	RW	20
Laurent Koscielny	Arsenal	31	CB	22

from www.transfermarkt.us

Example: Get the data

name	club	age	position	market value
Alexis Sanchez	Arsenal	28	LW	65
Mesut Ozil	Arsenal	28	AM	50
		5	GK	7
		3	RW	20
		1	CB	22

- Credible/Trustworthy?
- Possibly subjective market values?
- Sampled data

from www.transfermarkt.us

Example

name	club	age	position	market value
Alexis Sanchez	Arsenal	28	LW	65
Mesut Ozil	Arsenal	28	AM	50
Petr Cech	Arsenal	35	GK	7
Theo Walcott	Arsenal	28	RW	20
Laurent Koscielny	Arsenal	31	CB	22

Example: Explore the Data

name	club	age	position	market value
Alexis Sanchez	Arsenal	28	LW	65
Mesut Ozil	Arsenal	28	AM	50
Petr Cech	Arsenal	35	GK	7
Theo Walcott	Arsenal	28	RW	20
Laurent Koscielny				22

Does it contain the necessary information?

Example: Explore the Data

name	club	age	position	market value
Alexis Sanchez	Arsenal	28	LW	65
Mesut Ozil	Arsenal	28	AM	50
Petr Cech	Arsenal	35	GK	7
Theo Walcott	Arsenal	28	RW	20
Laurent Koscielny	Arsenal	31	CB	22

Missing data? Imputation needed?

Example: Explore the Data

name	club	age	position	market value
Alexis Sanchez	Arsenal	28	LW	65
Mesut Ozil	Arsenal	28	AM	50
Petr Cech	Arsenal	35	GK	7
Theo Walcott	Arsenal	28	RW	20
Laurent Koscielny	Arsenal	31	CB	22

Are the data types okay (`df.dtypes`)? Should be casted?

Example: Explore the Data

name	club	age	position	market value
Alexis Sanchez	Arsenal	28	LW	65
Mesut Ozil	Arsenal	28	AM	50
Petr Cech	Arsenal	35	GK	7
Theo Walcott	Arsenal	28	RW	20
Laurent Koscielny	Arsenal	31	CB	22

Are the values reasonable? `DataFrame.describe()` …

Example: Explore the Data

	age	page_views	fpl_value	fpl_points	market_value
count	461.000000	461.000000	461.000000	461.000000	461.000000
mean	26.804772	763.776573	5.447939	57.314534	11.012039
std	3.961892	931.805757	1.346695	53.113811	12.257403
min	17.000000	3.000000	4.000000	0.000000	0.050000
25%	24.000000	220.000000	4.500000	5.000000	3.000000
50%	27.000000	460.000000	5.000000	51.000000	7.000000
75%	30.000000	896.000000	5.500000	94.000000	15.000000
max	38.000000	7664.000000	12.500000	264.000000	75.000000

Are the values reasonable? `DataFrame.describe()` …

Example: Explore the Data

	age	page_views	fpl_value	fpl_points	market_value
count	461.000000	461.000000	461.000000	461.000000	461.000000
mean	26.804772	763.776573	5.447939	57.314534	11.012039
std	3.961892	931.805757	1.346695	53.113811	12.257403
min	17.000000	3.000000	4.000000	0.000000	0.050000
25%	24.000000	220.000000	4.500000	5.000000	3.000000
50%	27.000000	460.000000	5.000000	51.000000	7.000000
75%	30.000000	896.000000	5.500000	94.000000	15.000000
max	38.000000	7664.000000	12.500000	264.000000	75.000000

Summary statistics can only reveal so much

Lecture Outline

Data Science Process Example

- Dataset considerations
  - Comprehensive vs Sampled
  - Biases

Visualization

- Exploration (EDA)
- Communication

Visualization

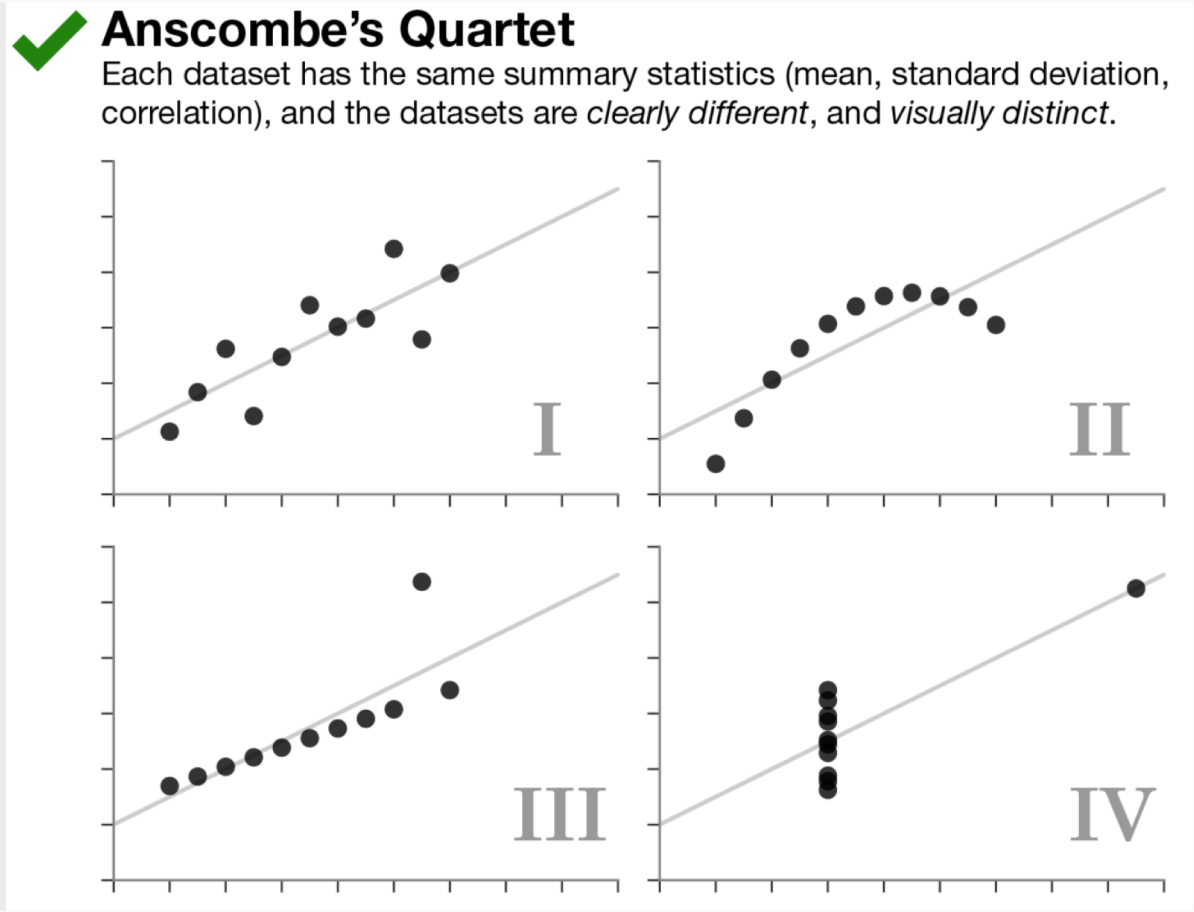

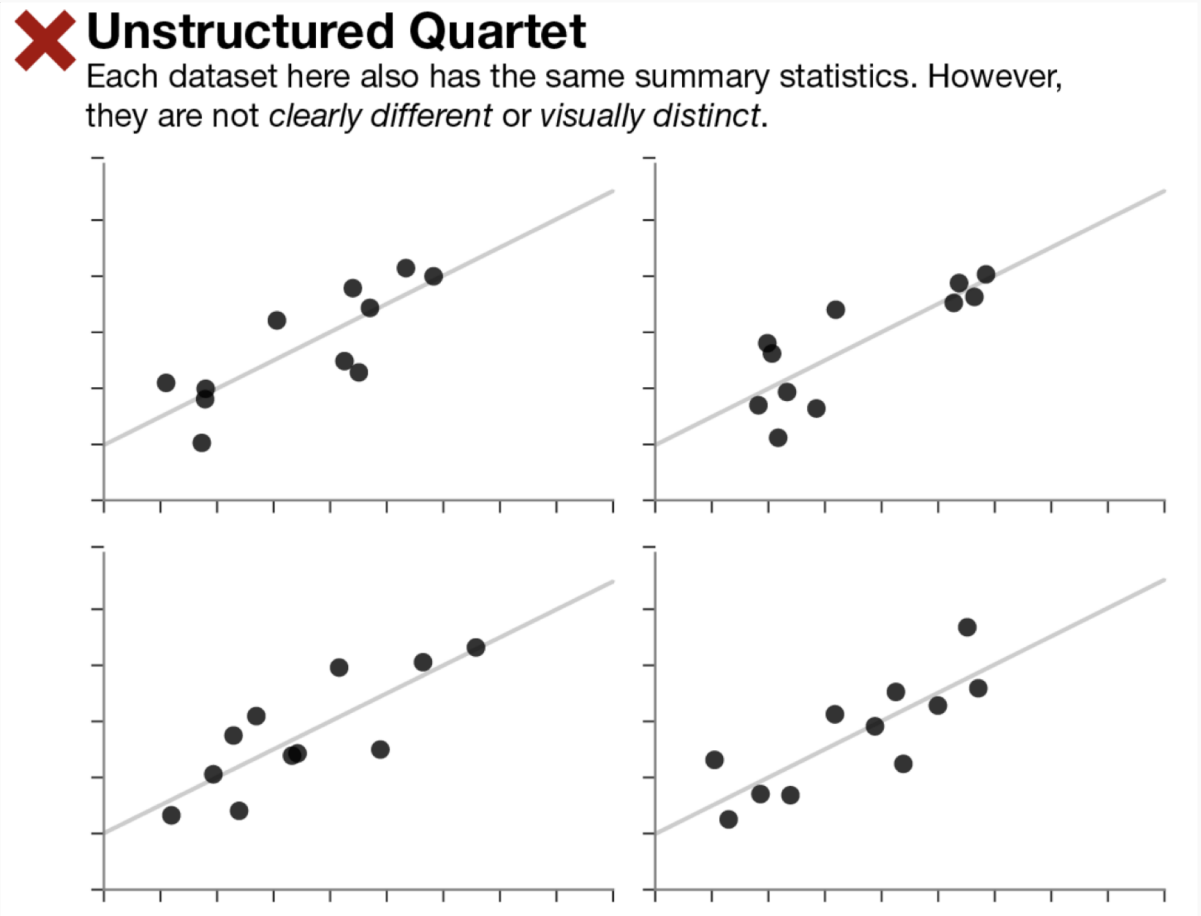

Same stats do not imply same graphs

Same graphs do not imply same stats

Visualization

	Antibiotic			
Bacteria	Penicillin	Streptomycin	Neomycin	Gram Staining
Aerobacter aerogenes	870	1	1.6	negative
Brucella abortus	1	2	0.02	negative
Brucella anthracis	0.001	0.01	0.007	positive
Diplococcus pneumoniae	0.005	11	10	positive
Escherichia coli	100	0.4	0.1	negative
Klebsiella pneumoniae	850	1.2	1	negative
Mycobacterium tuberculosis	800	5	2	negative
Proteus vulgaris	3	0.1	0.1	negative
Pseudomonas aeruginosa	850	2	0.4	negative
Salmonella (Eberthella) typhosa	1	0.4	0.008	negative
Salmonella schottmuelleri	10	0.8	0.09	negative
Staphylococcus albus	0.007	0.1	0.001	positive
Staphylococcus aureus	0.03	0.03	0.001	positive
Streptococcus fecalis	1	1	0.1	positive
Streptococcus hemolyticus	0.001	14	10	positive
Streptococcus viridans	0.005	10	40	positive

Visualization

	Antibiotic			
Bacteria	Penicillin	Streptomycin	Neomycin	Gram Staining
Aerobacter aerogenes	870	1	1.6	negative
Brucella abortus	1	2	0.02	negative
Brucella anthracis	0.001	0.01	0.007	positive
Diplococcus pneumoniae	0.005	11	10	positive
Escherichia coli	100	0.4	0.1	negative
Klebsiella pneumoniae	850	1.2	1	negative
Mycobacterium tuberculosis	800	5	2	negative
Proteus vulgaris	3	0.1	0.1	negative
Pseudomonas aeruginosa	850	2	0.4	negative
Salmonella (Eberthella) typhosa	1	0.4	0.008	negative
Salmonella schottmuelleri	10	0.8	0.09	negative
Staphylococcus albus	0.007	0.1	0.001	positive
Staphylococcus aureus	0.03	0.03	0.001	positive
Streptococcus fecalis	1	1	0.1	positive
Streptococcus hemolyticus	0.001	14	10	positive
Streptococcus viridans	0.005	10	40	positive

Visualization

	Antibiotic			
Bacteria	Penicillin	Streptomycin	Neomycin	Gram Staining
Aerobacter aerogenes	870	1	1.6	negative
Brucella abortus	1	2	0.02	negative
Brucella anthracis	0.001	0.01	0.007	positive
Diplococcus pneumoniae	0.005	11	10	positive
Escherichia coli	100	0.4	0.1	negative
Klebsiella pneumoniae	850	1.2	1	negative
Mycobacterium tuberculosis	800	5	2	negative
Proteus vulgaris	3	0.1	0.1	negative
Pseudomonas aeruginosa	850	2	0.4	negative
Salmonella (Eberthella) typhosa	1	0.4	0.008	negative
Salmonella schottmuelleri	10	0.8	0.09	negative
Staphylococcus albus	0.007	0.1	0.001	positive
Staphylococcus aureus	0.03	0.03	0.001	positive
Streptococcus viridans	0.005	10	40	positive

What are some questions we could ask?

Visualization

	Antibiotic			
Bacteria	Penicillin	Streptomycin	Neomycin	Gram Staining
Aerobacter aerogenes	870	1	1.6	negative
Brucella abortus	1	2	0.02	negative
Brucella anthracis	0.001	0.01	0.007	positive
Diplococcus pneumoniae	0.005	11	10	positive
Escherichia coli	100	0.4	0.1	negative
Klebsiella pneumoniae	850	1.2	1	negative
Mycobacterium tuberculosis	800	5	2	negative
Proteus vulgaris	3	0.1	0.1	negative
Pseudomonas aeruginosa	850	2	0.4	negative
Salmonella (Eberthella) typhosa	1	0.4	0.008	negative
Salmonella schottmuelleri	10	0.8	0.09	negative
Staphylococcus albus	0.007	0.1	0.001	positive
Staphylococcus aureus	0.03	0.03	0.001	positive
Streptococcus viridans	0.005	10	40	positive

Q: How effective are the antibiotics?

Gram Positive

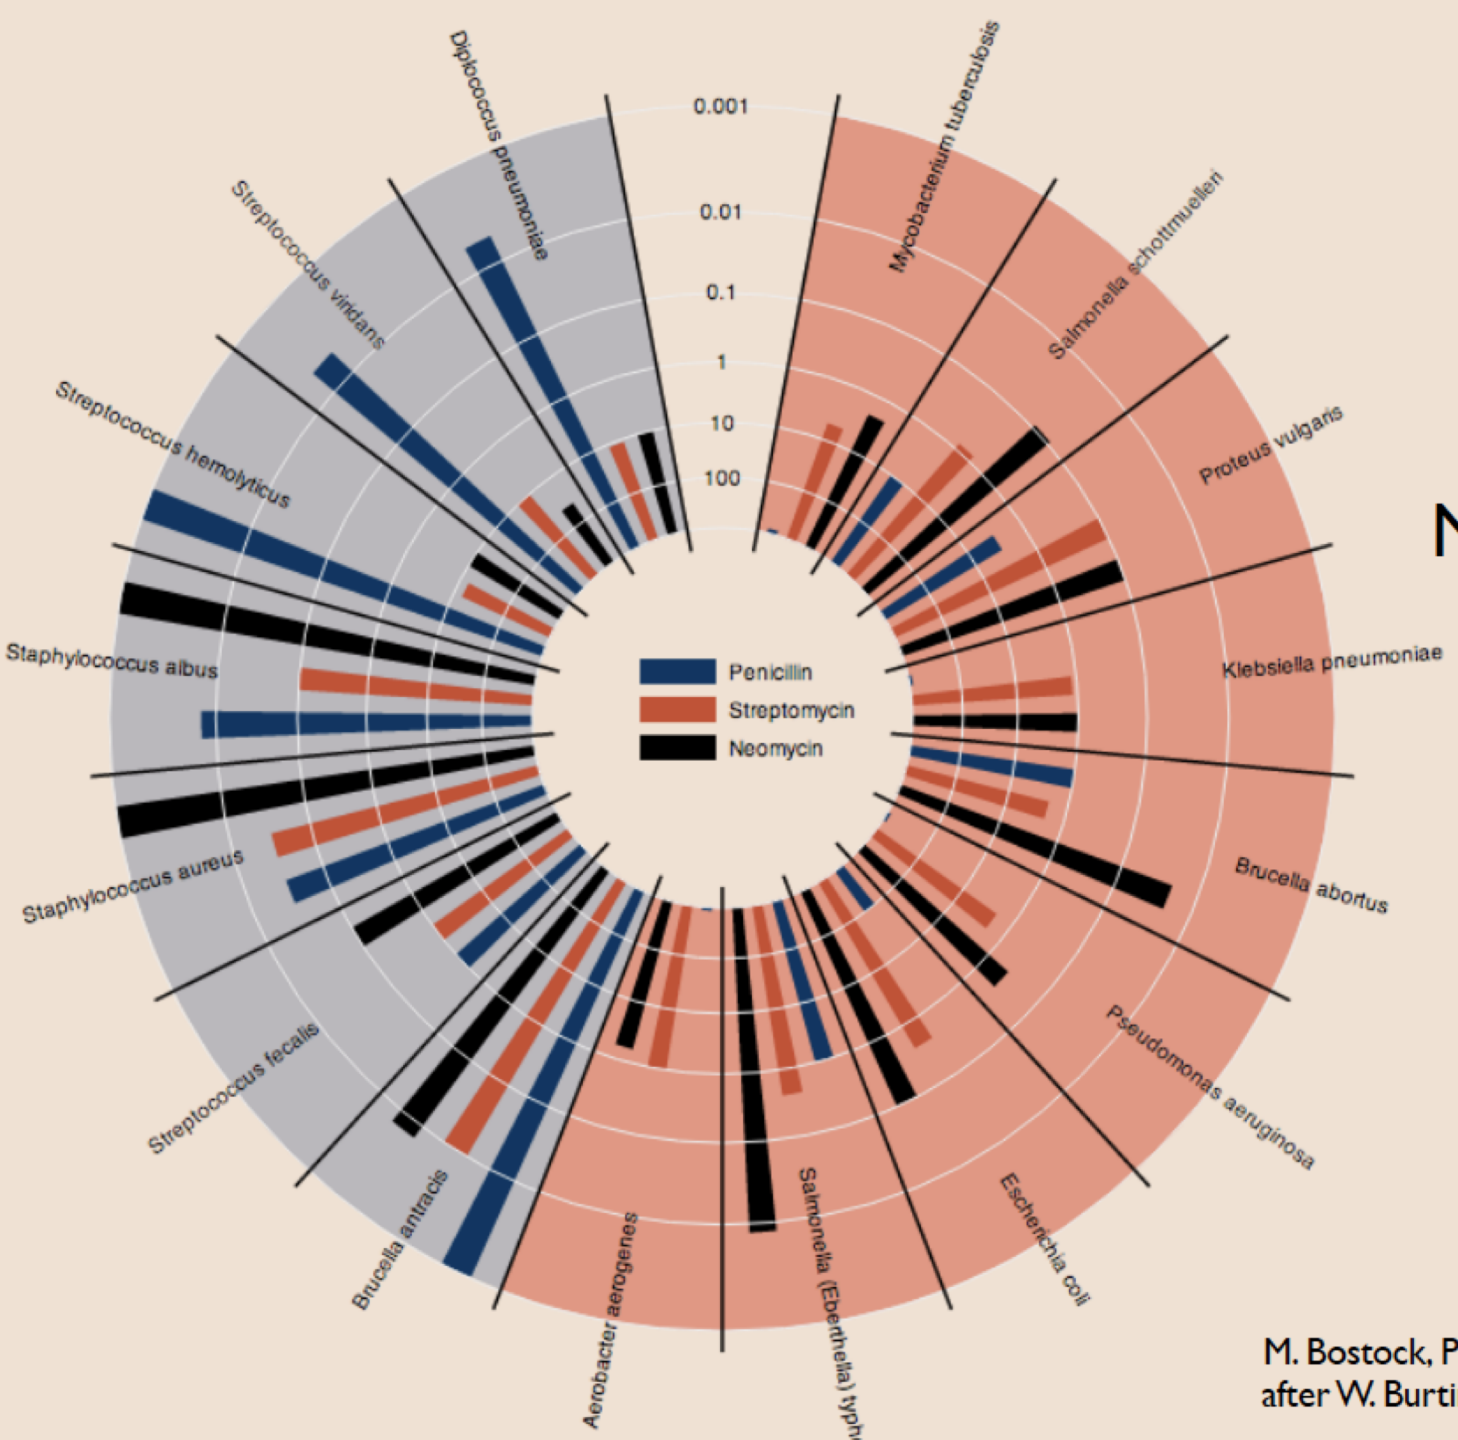

Gram Negative

M. Bostock, Protovis
after W. Burtin, 1951

Gram Positive

Gram Negative

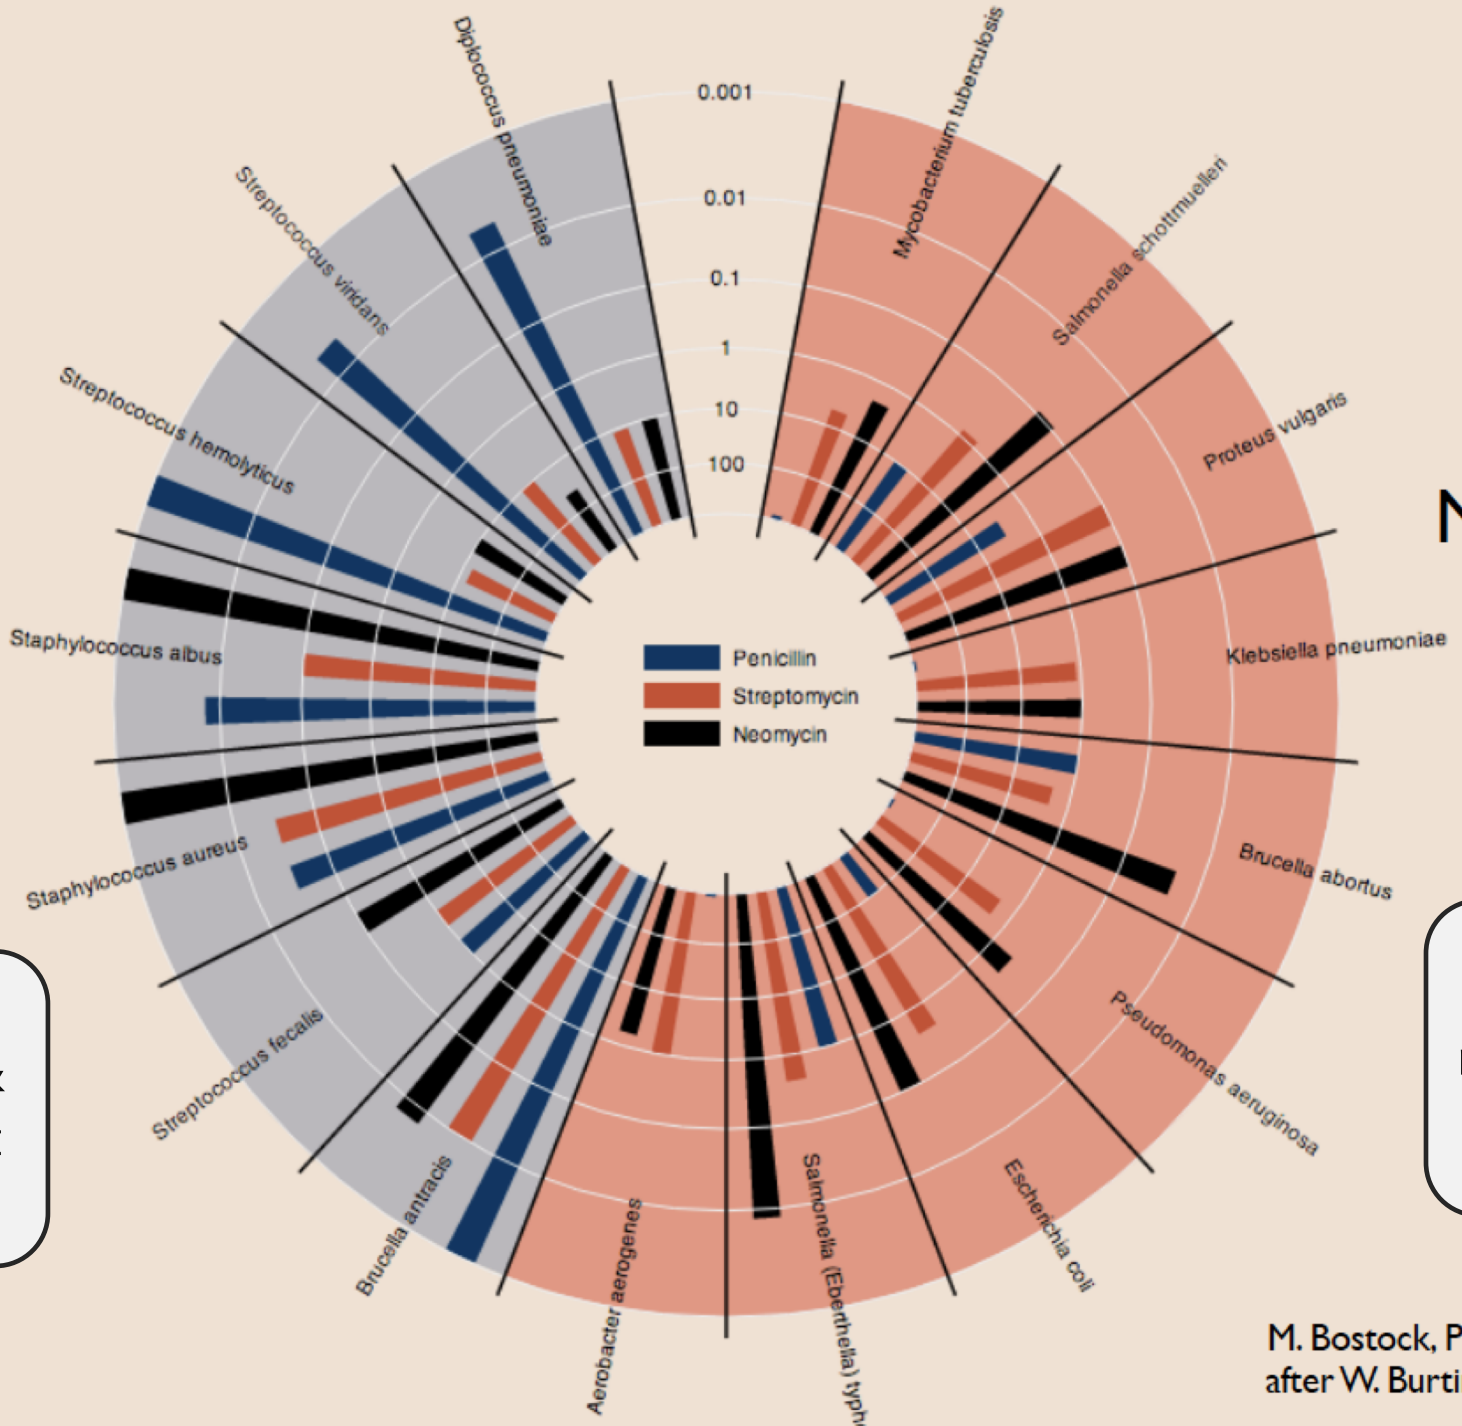

If bacteria is gram positive, Penicillin & Neomycin are most effective

If bacteria is gram negative, Neomycin is most effective

M. Bostock, Protovis after W. Burtin, 1951

How do the bacteria compare?

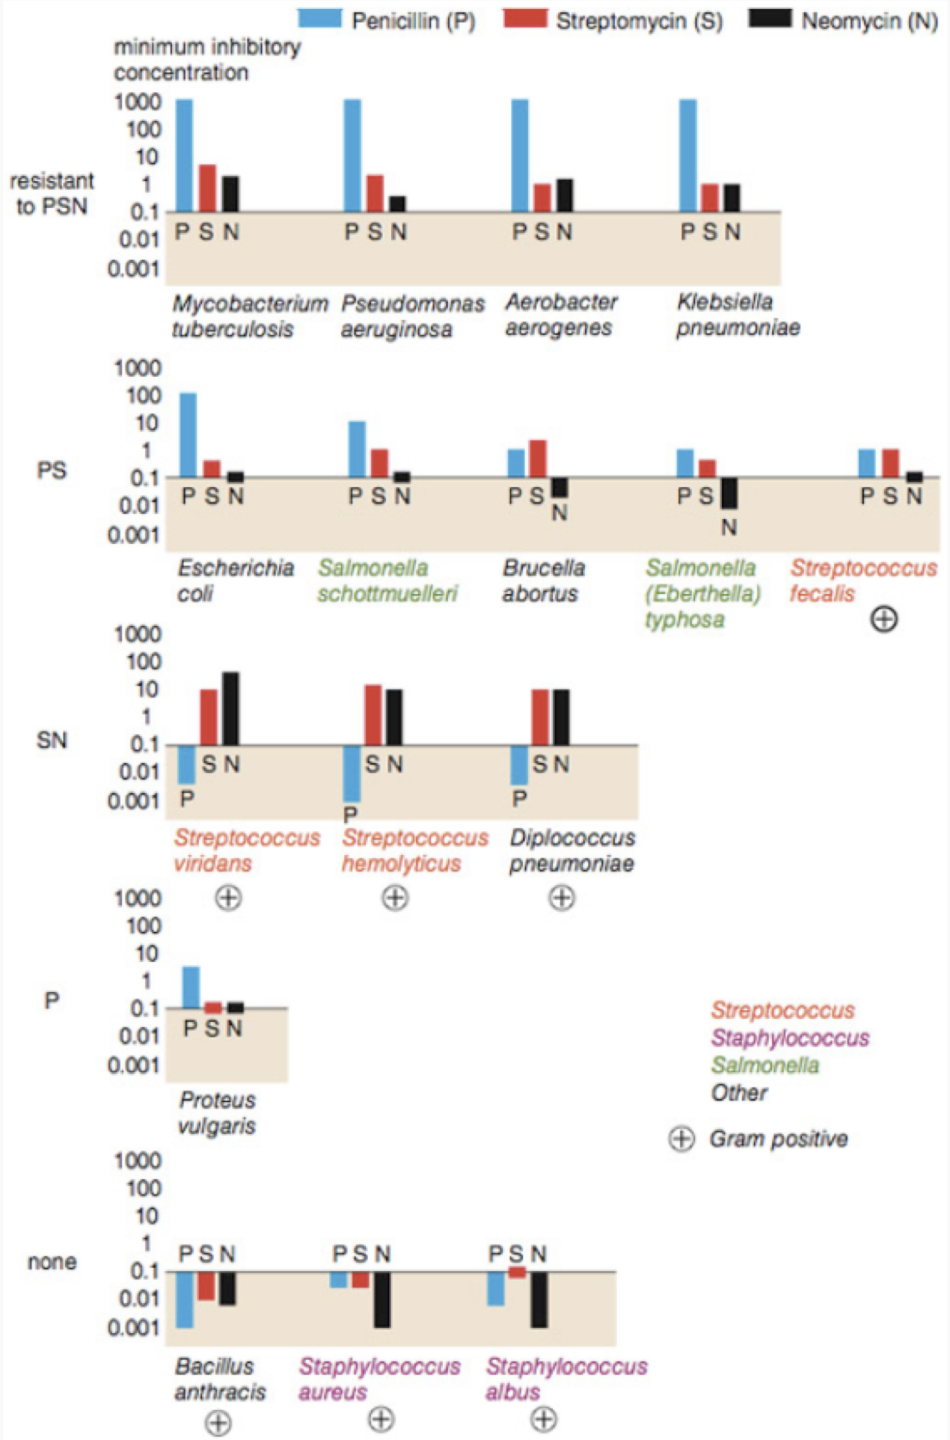

Wainer & Lysen, "That's funny..."
American Scientist, 2009
Adapted from Brian Schmotzer

How do the bacteria compare?

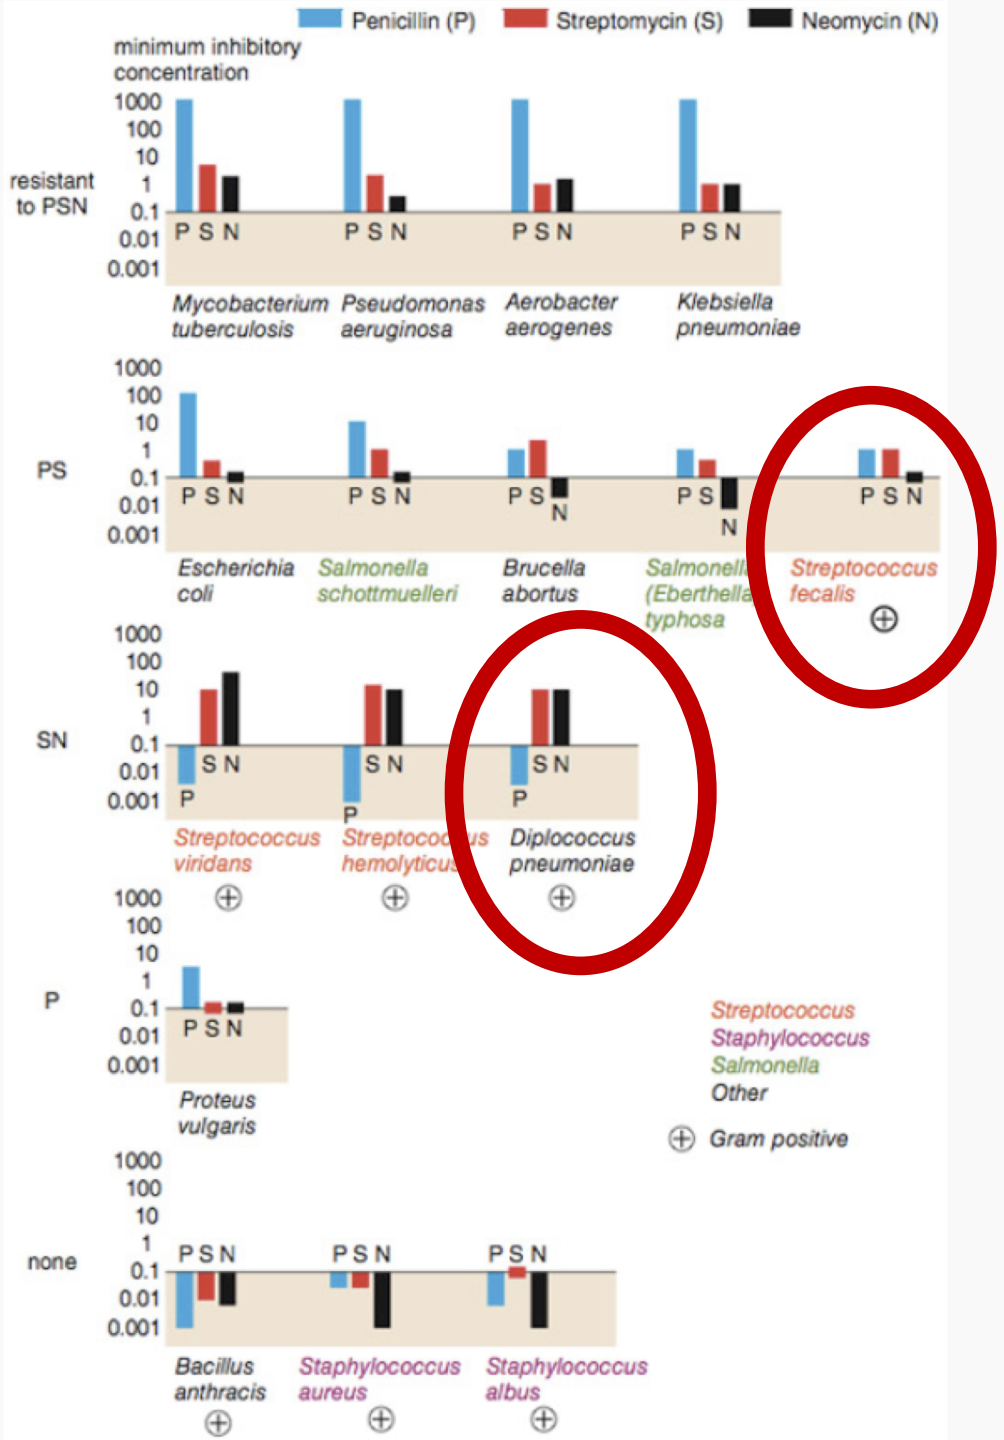

Wainer & Lysen, "That's funny..."
American Scientist, 2009
Adapted from Brian Schmotzer

How do the bacteria compare?

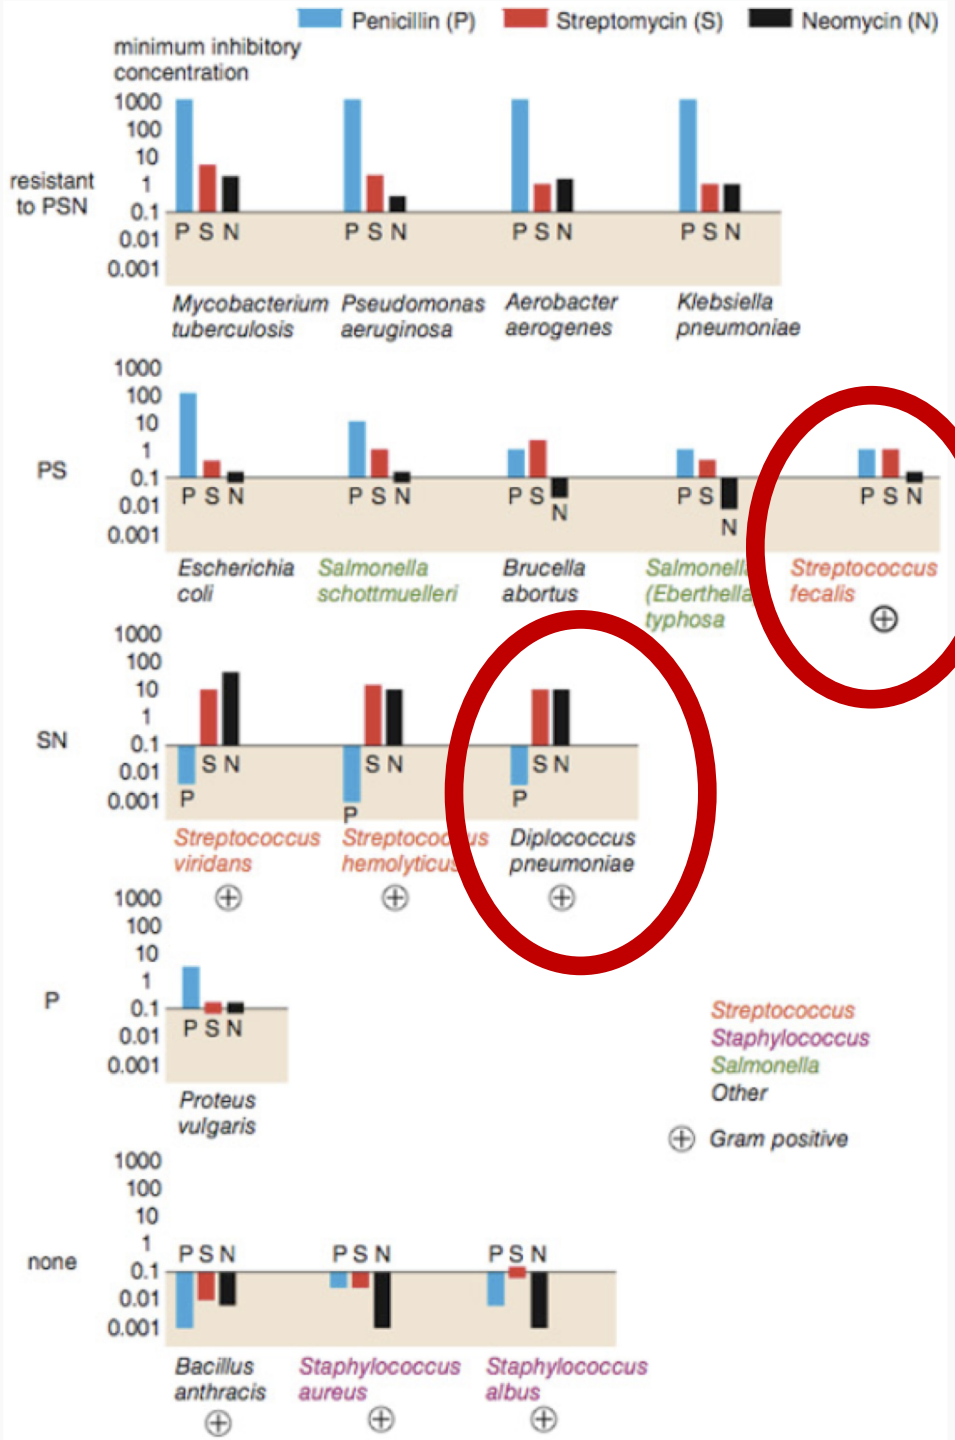

Not a streptococcus!
(realized ~30 years later)

Actually a streptococcus!
(realized ~20 years later)

Wainer & Lysen, "That's funny..."
American Scientist, 2009
Adapted from Brian Schmotzer

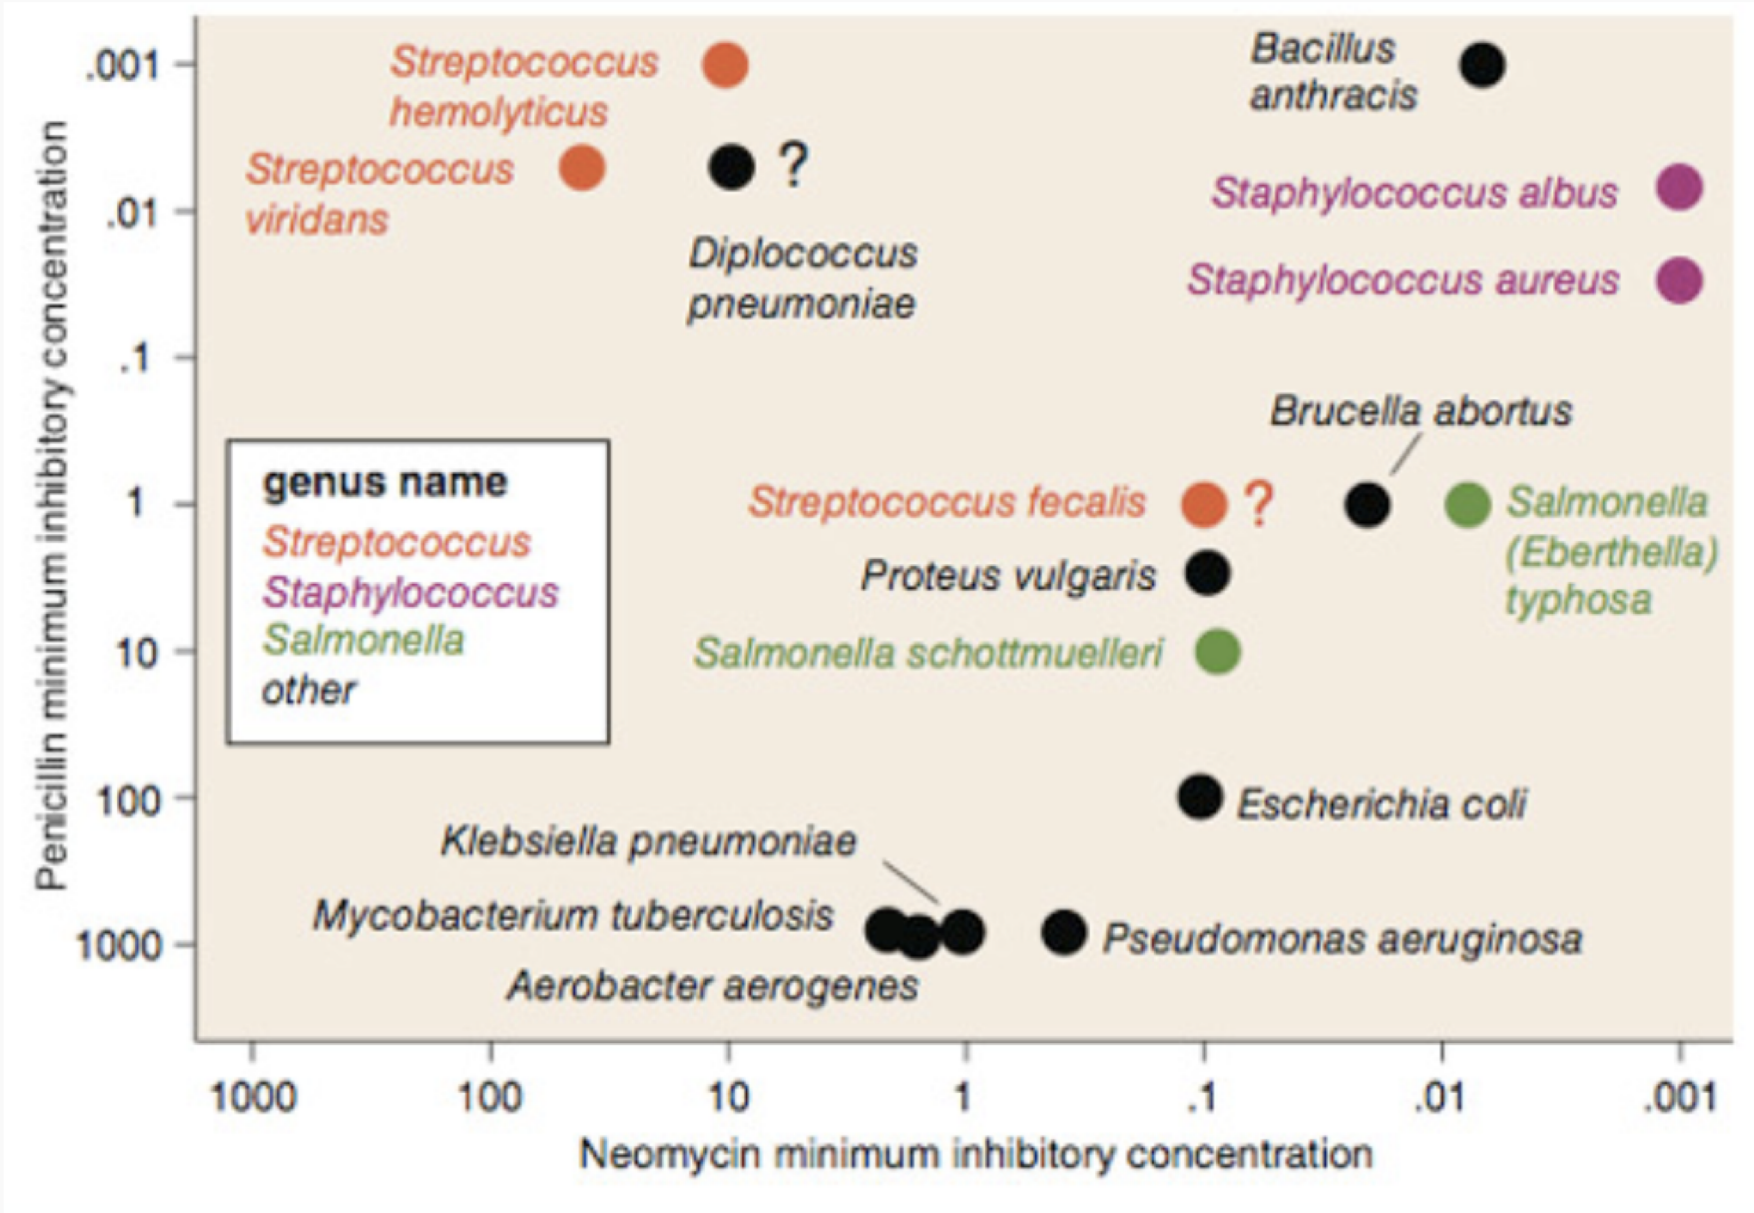

Wainer & Lysen, "That's funny..."
American Scientist, 2009

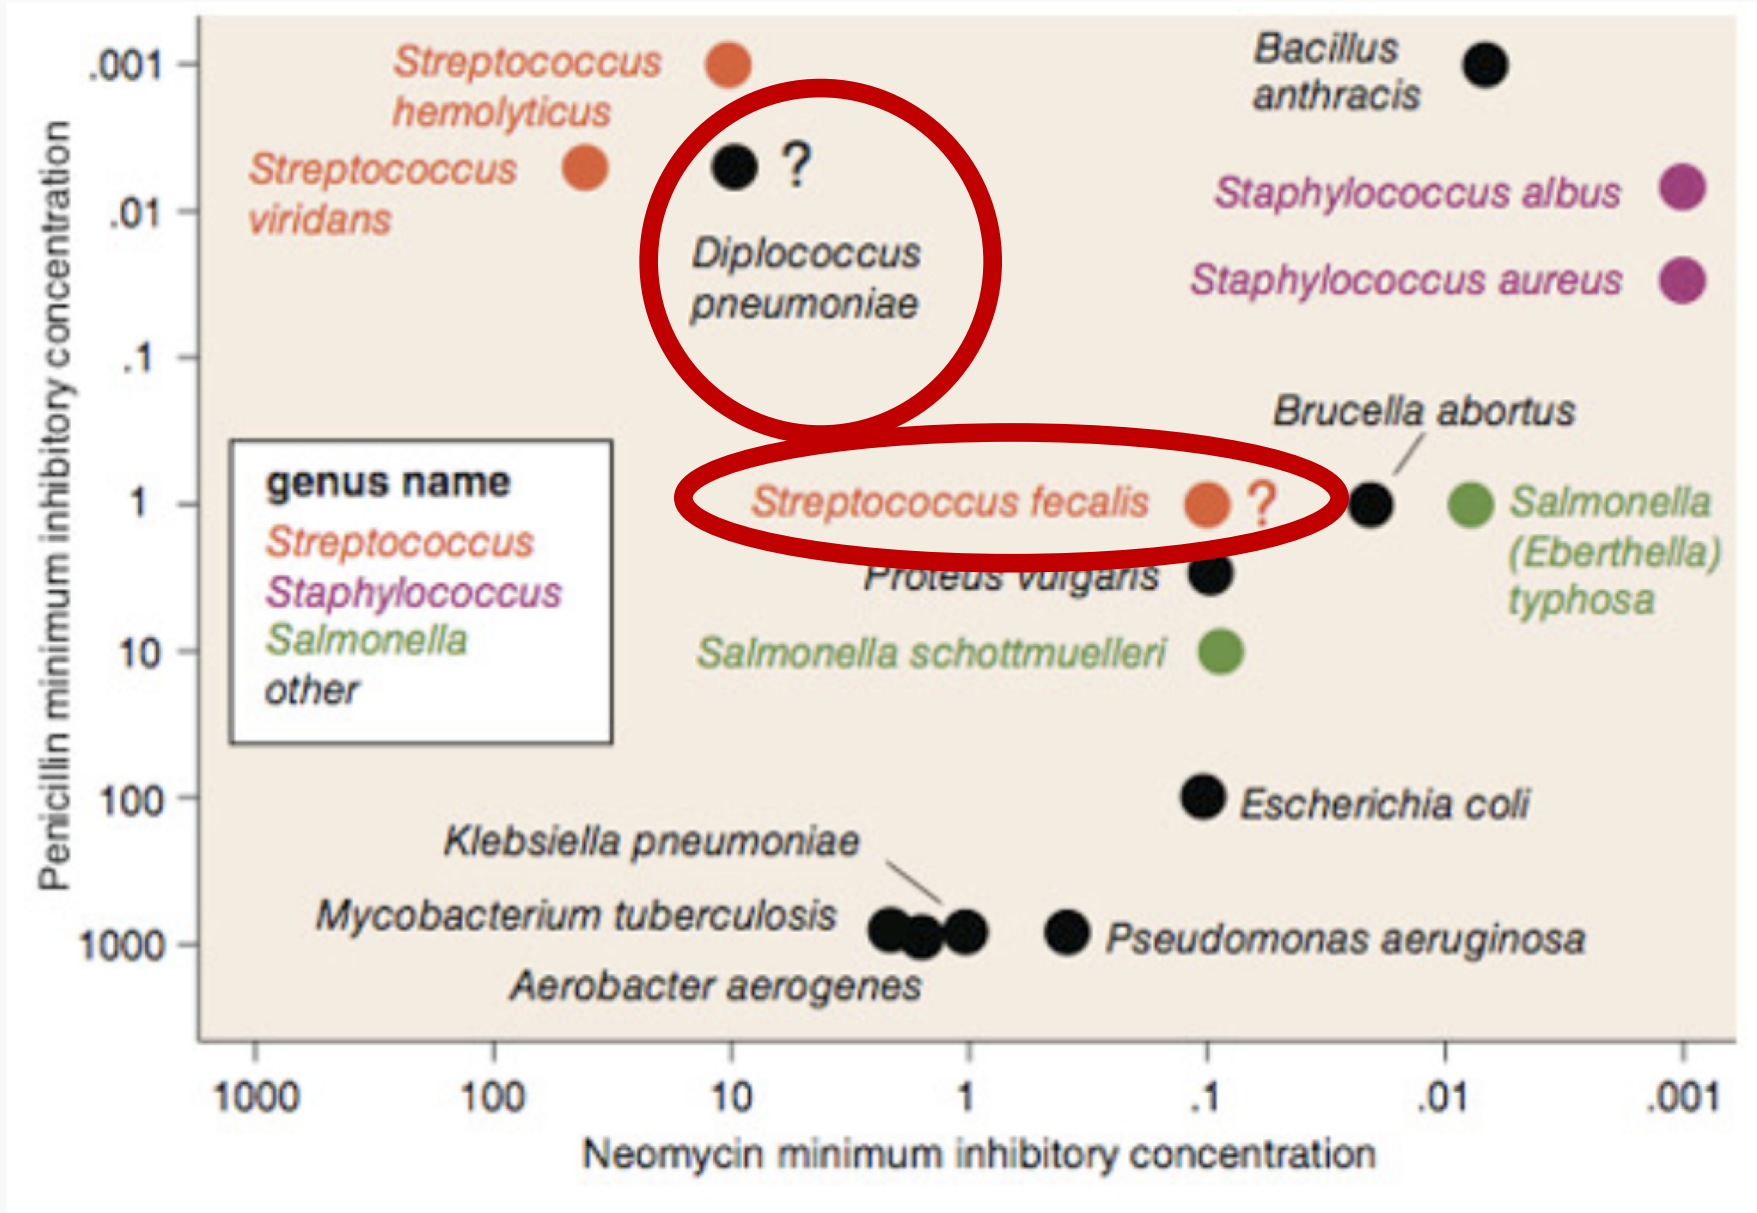

Wainer & Lysen, "That's funny..."
American Scientist, 2009

Visualization

"The greatest value of a picture is when it forces us to notice what we never expected to see."

John Tukey

Visualization Goals

Communicate (explanatory)

- Present data and ideas
- Explain and inform
- Provide evidence and support
- Influence and persuade

Analyze (exploratory)

- Explore the data
- Assess a situation
- Determine how to proceed
- Decide what to do

Visualization Goals

Communicate (explanatory)

- Present data and ideas
- Explain and inform
- Provide evidence and support
- Influence and persuade

Analyze (exploratory)

- Explore the data
- Assess a situation
- Determine how to proceed
- Decide what to do

You're essentially communicating drafts to yourself

Communicate

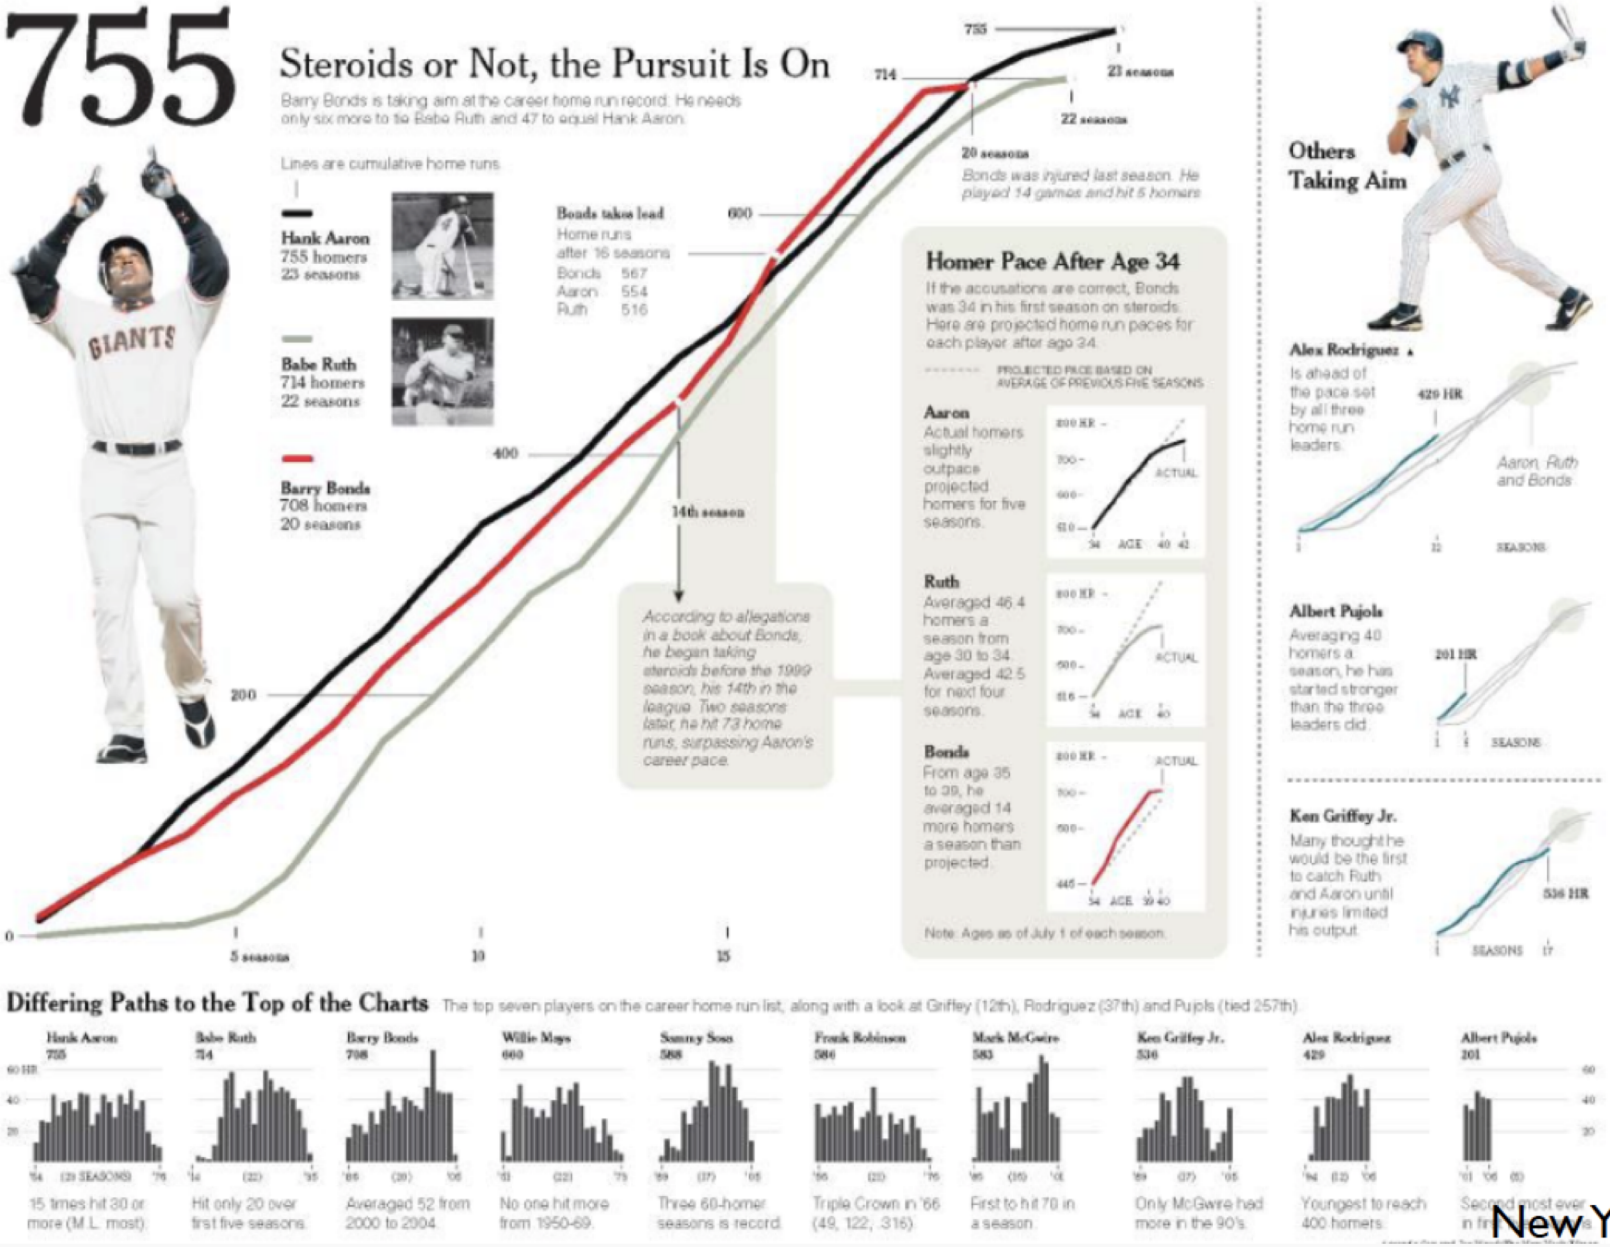

New York Times

Explore

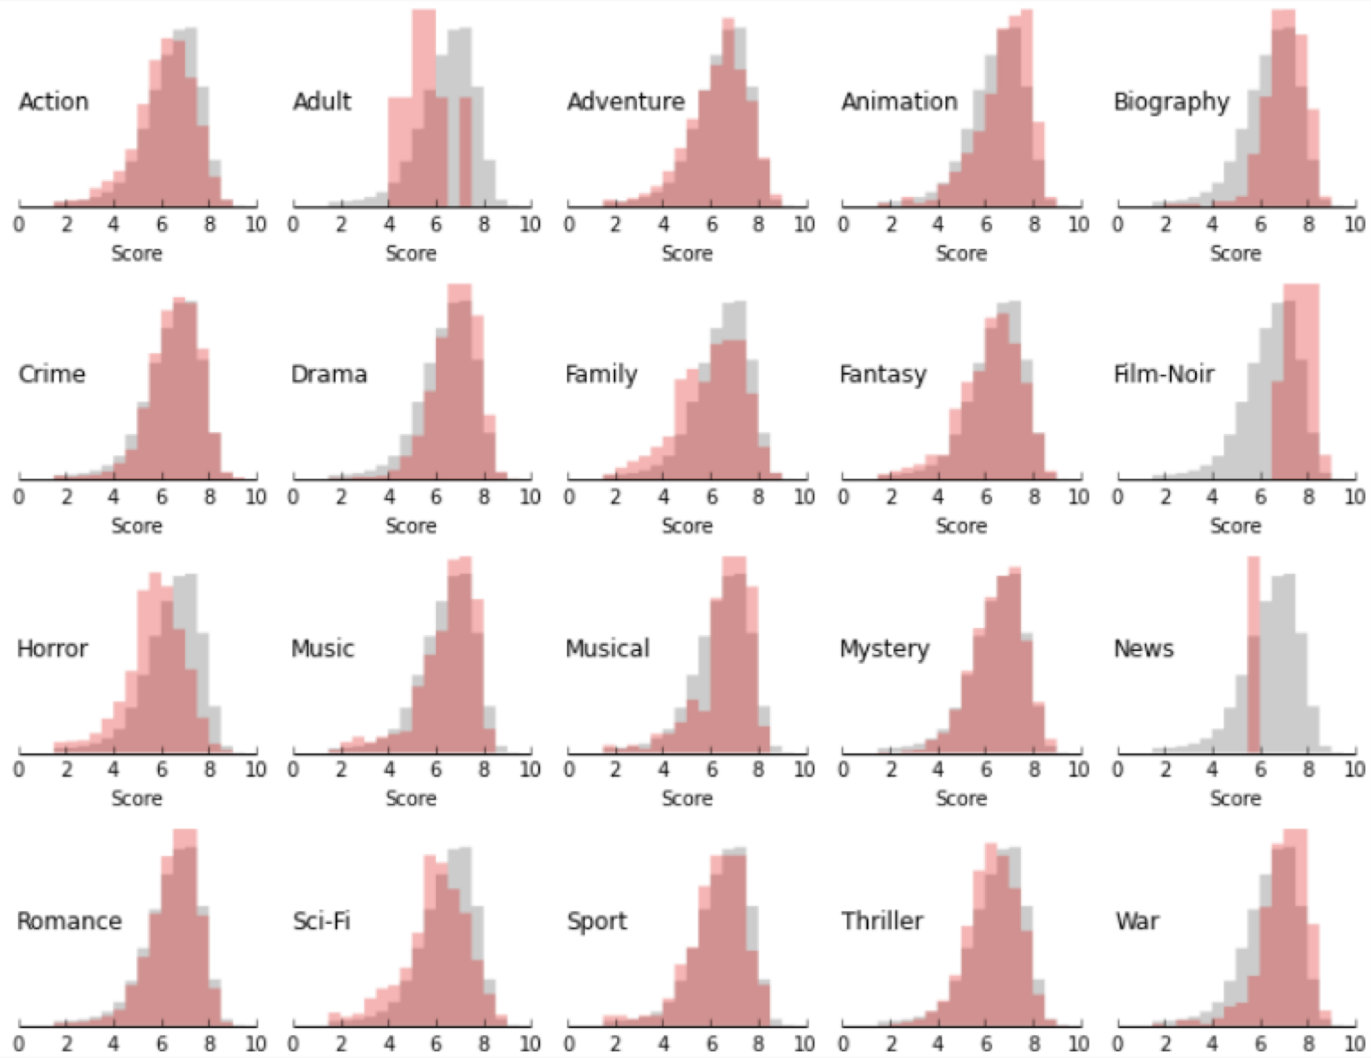
Action
Adult
Adventure
Animation
Biography
Crime
Drama
Family
Fantasy
Film-Noir
Horror
Music
Musical
Mystery
News
Romance
Sci-Fi
Sport
Thriller
War
0 2 4 6 8 10
Score

Not Effective

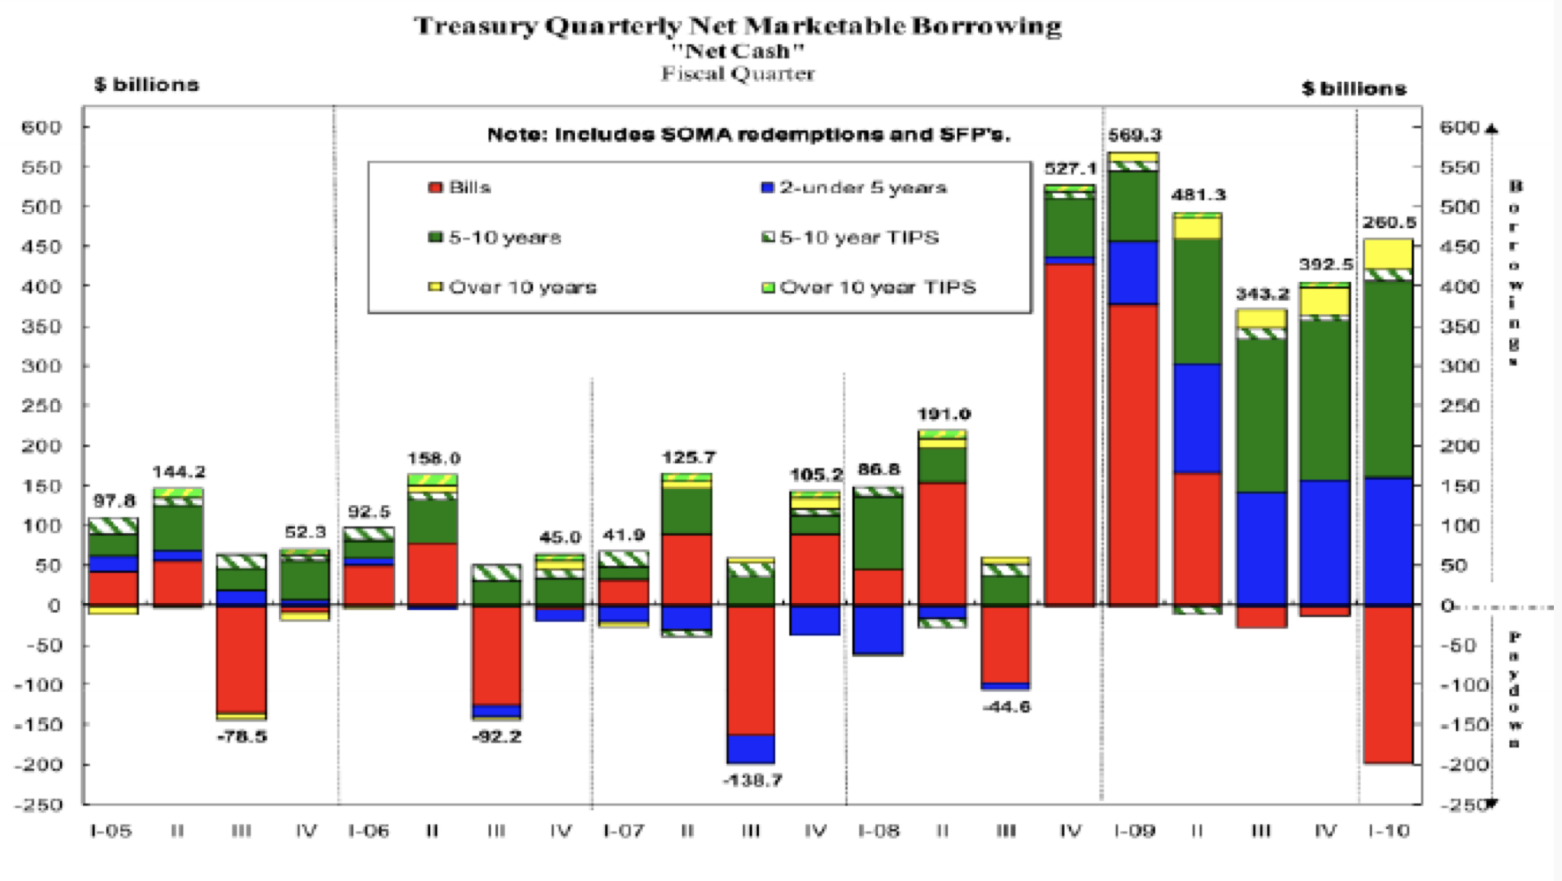

Not Effective

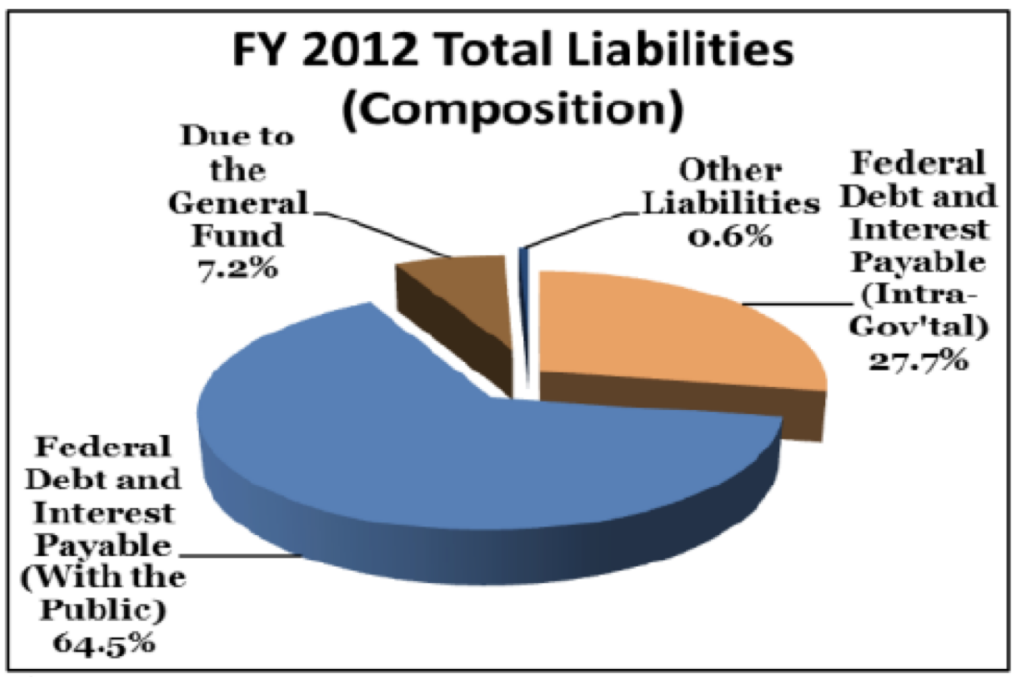

Figure 10

Not Effective

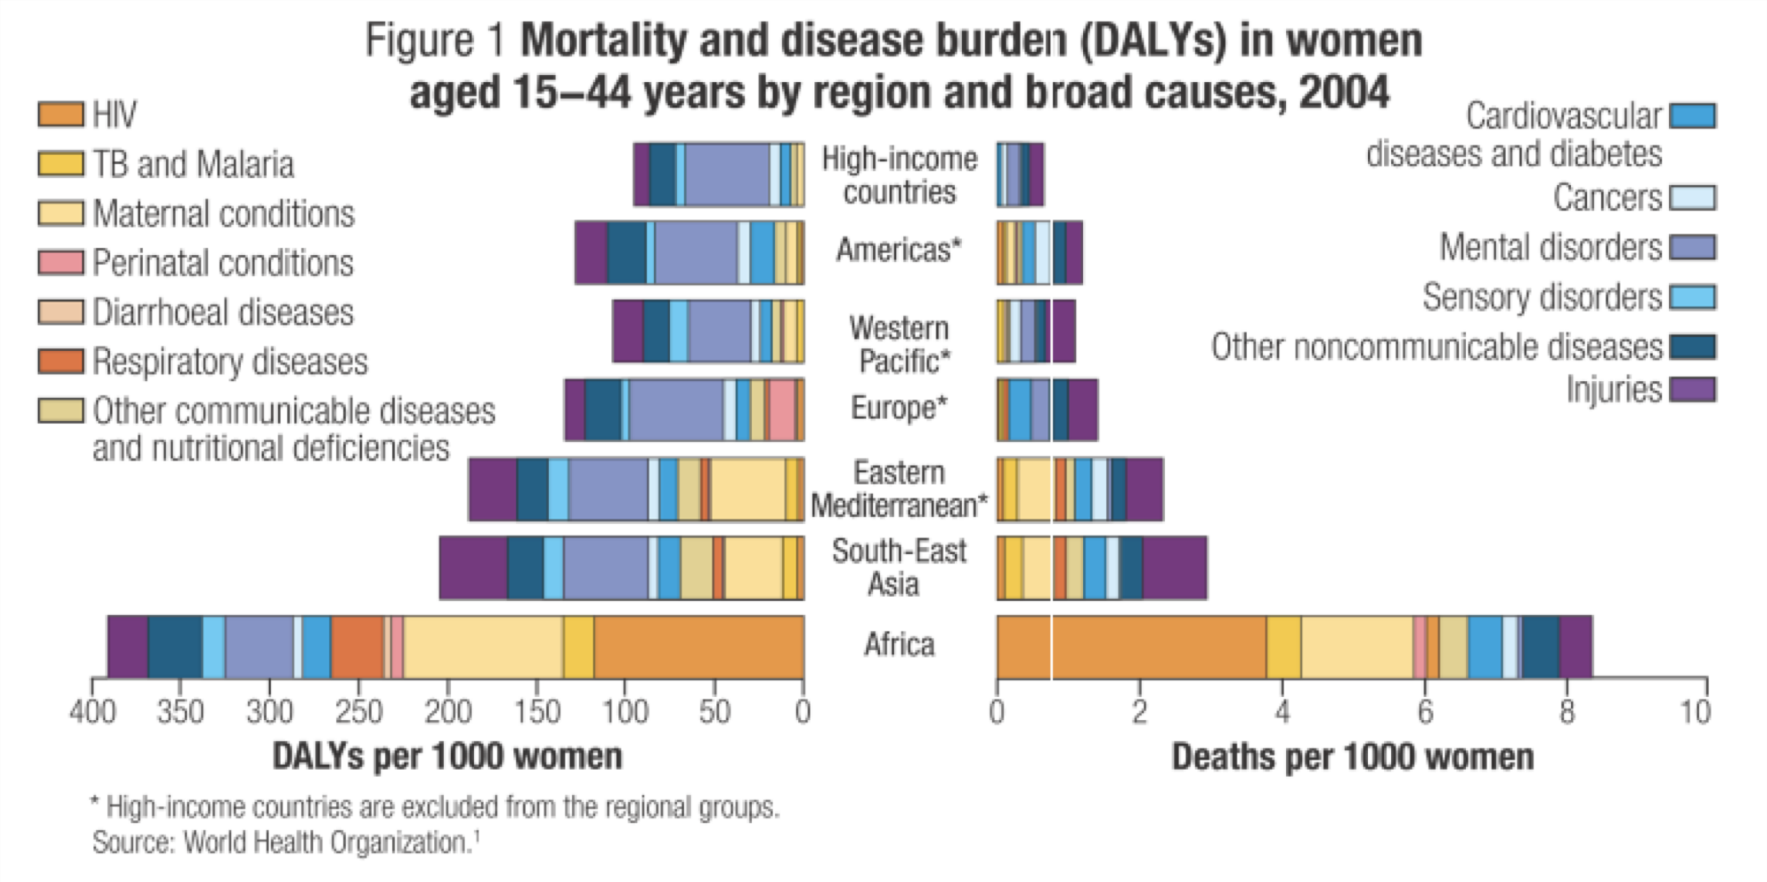

Not Effective

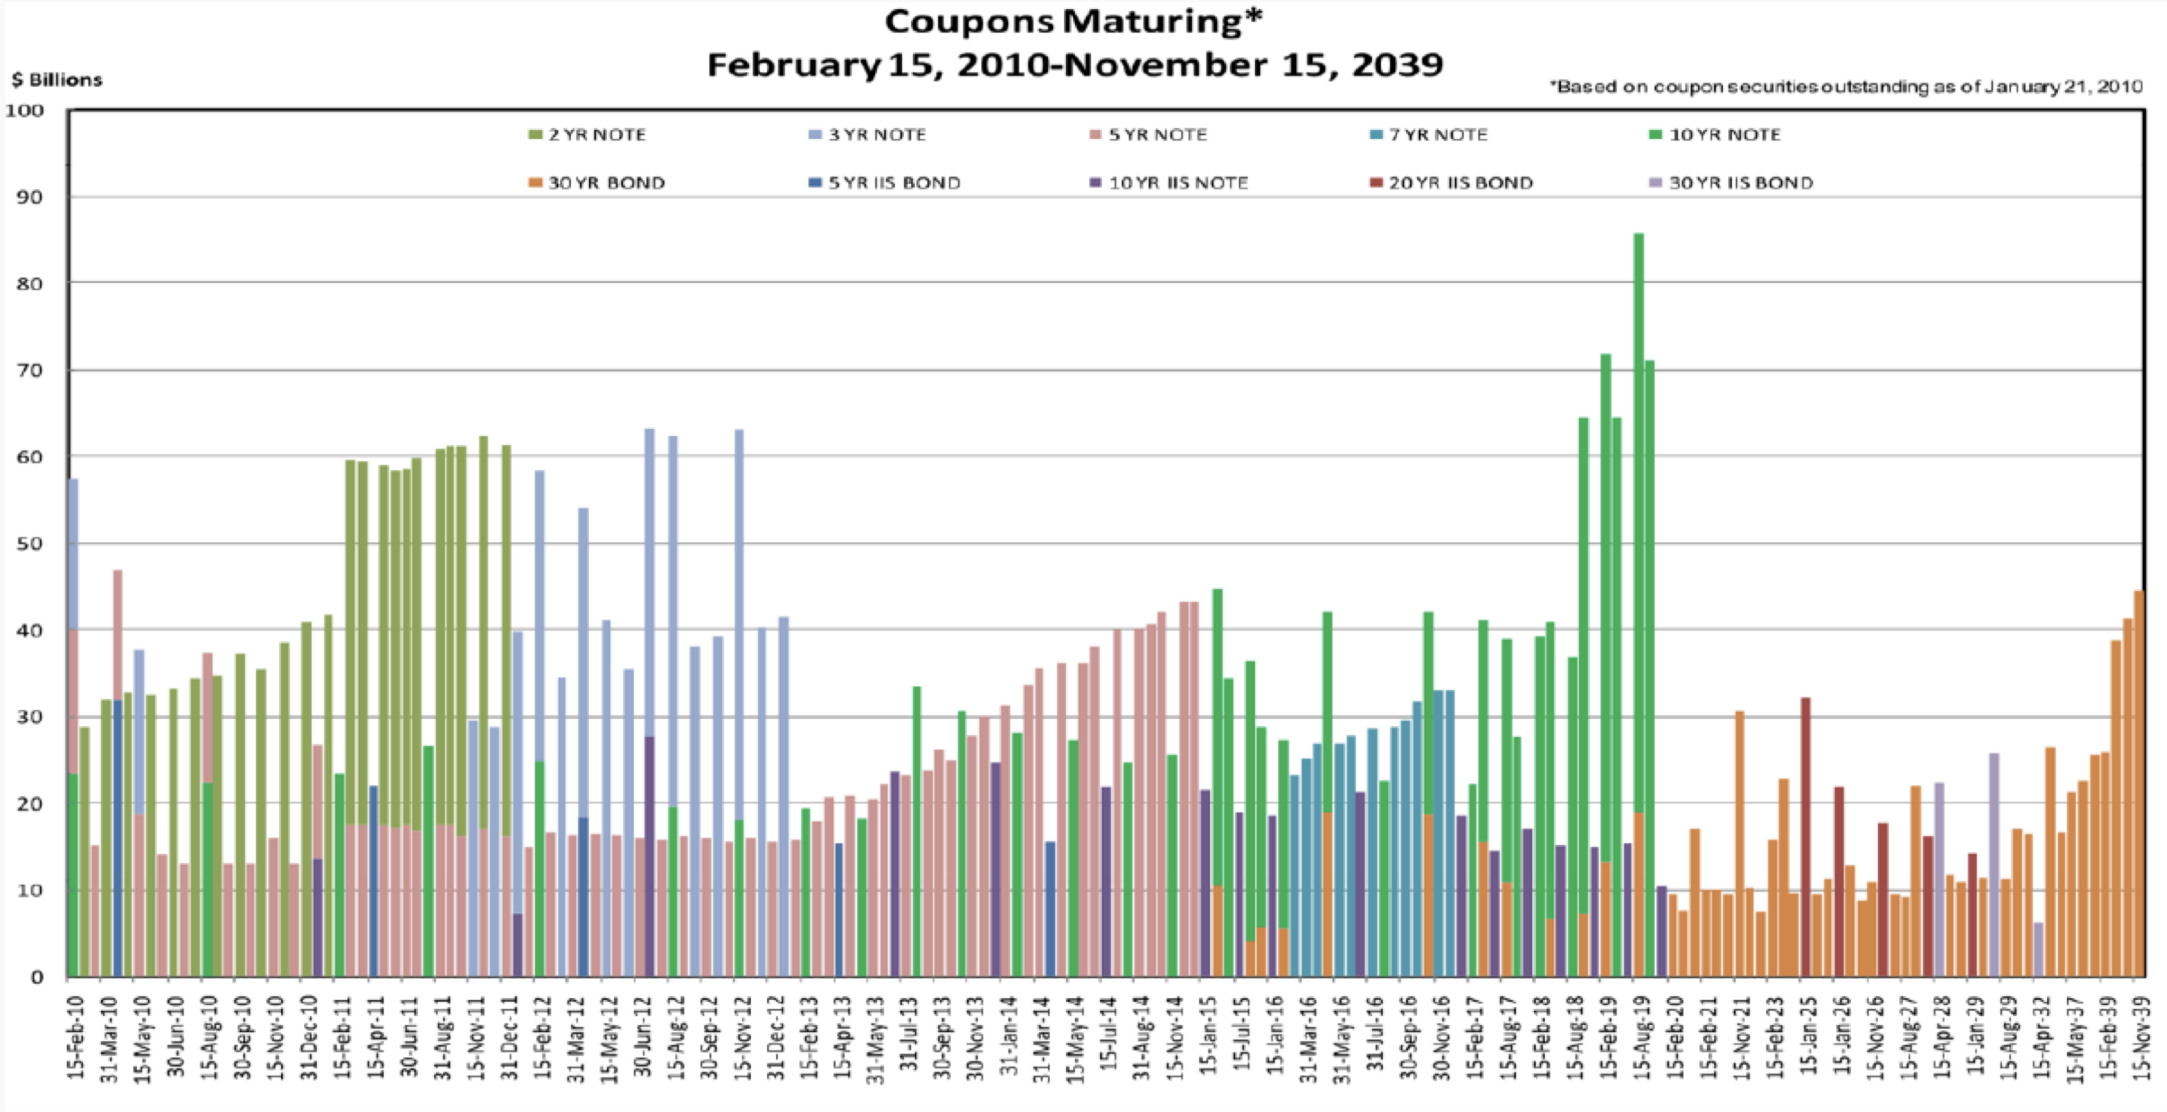

Not Effective

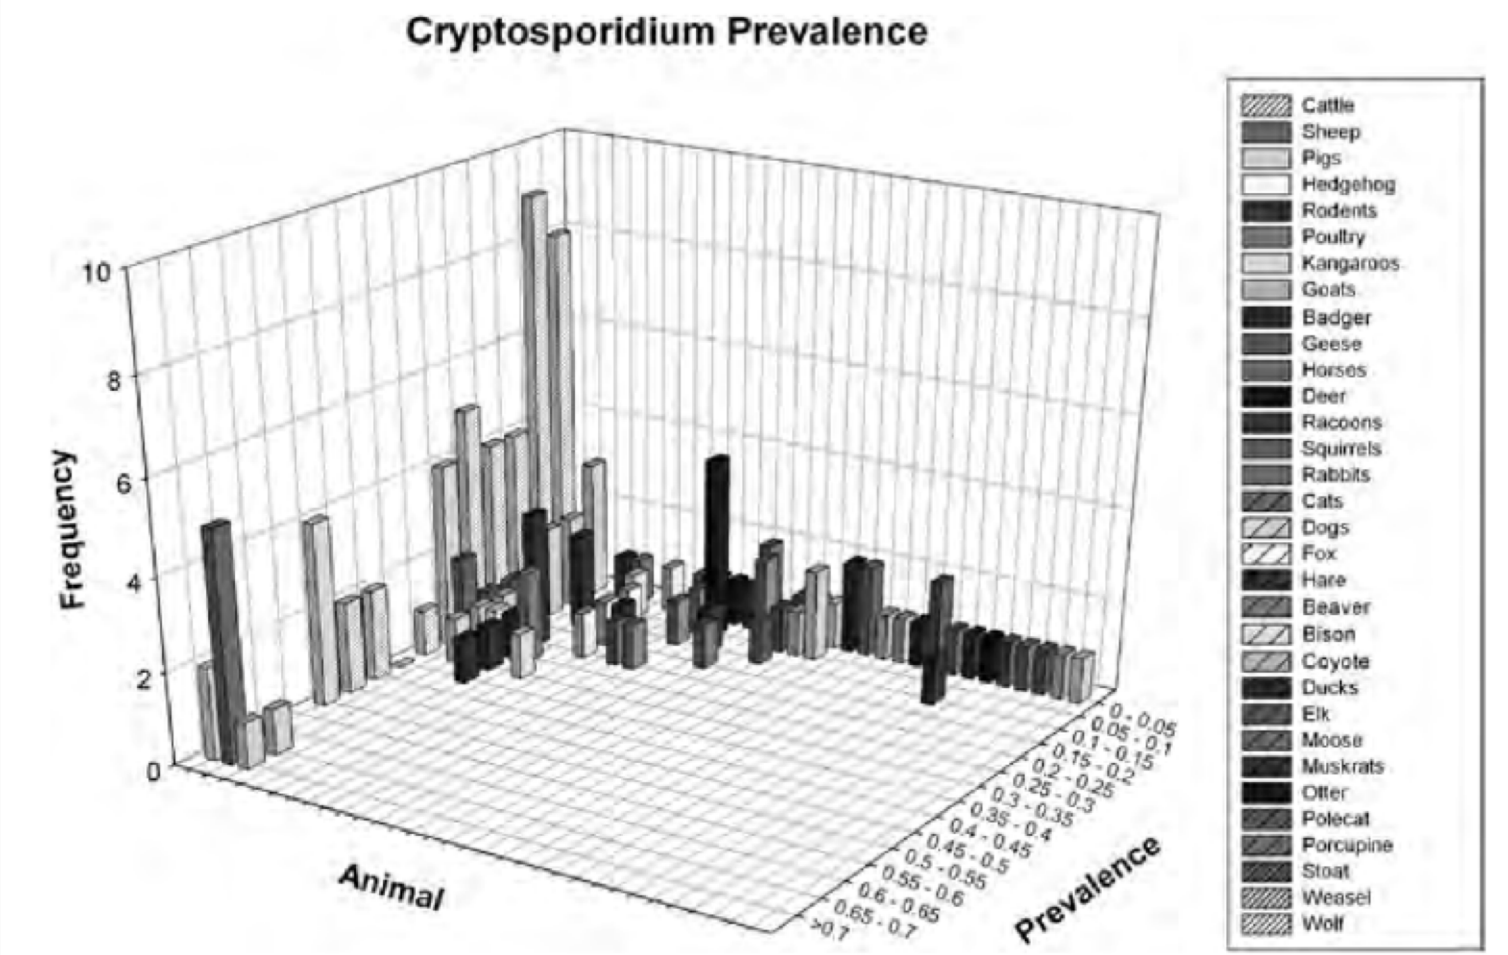

Figure 5.2 Mean prevalence rates of *Cryptosporidium* oocysts by animal species.

Visualization

Let's say that we are interested in the English Premier League (football/soccer) and want to build a model to predict a player's market value.

Question Does age affect one's market value?

What type of visualization would help us explore this question?